

LangChain is a framework for building applications powered by large language models (LLMs).

It enables developers to chain prompts, models, and memory for more complex workflows.

Gemini is a family of large language models developed by Google DeepMind.

It is designed for tasks such as reasoning, summarization, and answering complex questions.

Retrieval-Augmented Generation (RAG) improves accuracy by grounding LLM responses in external documents.